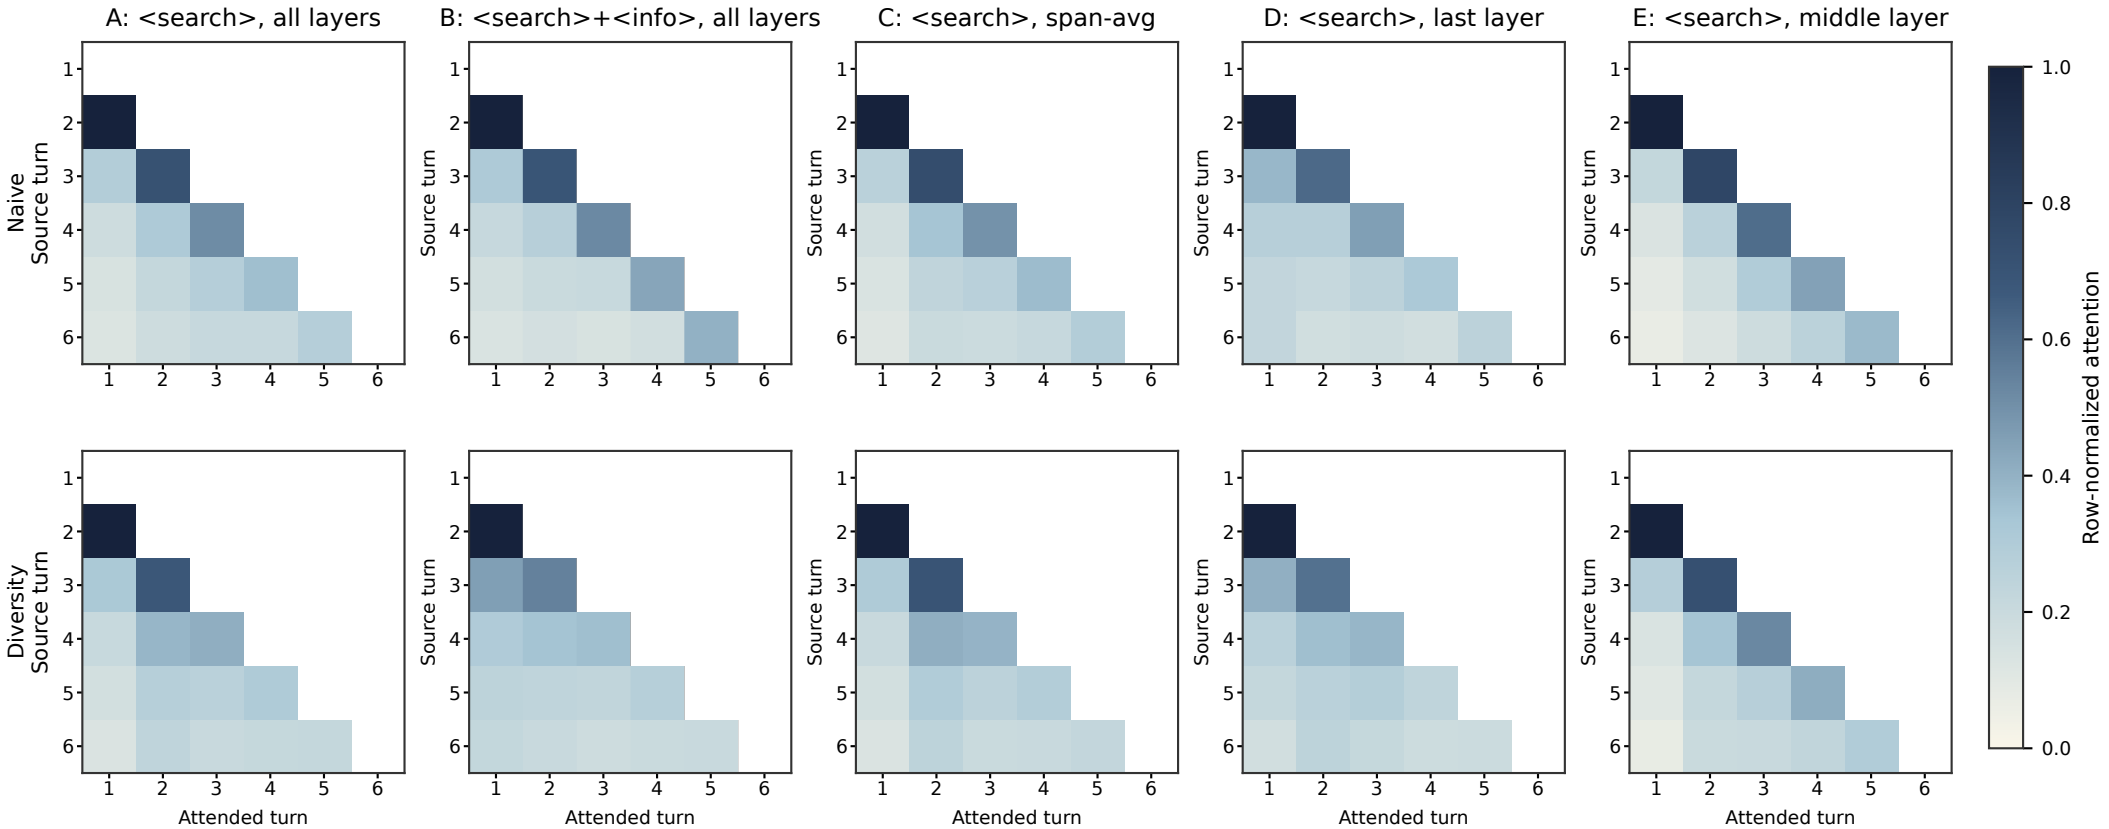
Attention bin variant comparison — qwen3-8b GAIA Q=gaia-97
A: <search>, all layers
B: <search>+<info>, all layers
C: <search>, span-avg
D: <search>, last layer
E: <search>, middle layer
Naive
Diversity
Source turn
Attended turn
Row-normalized attention
0.0
0.2
0.4
0.6
0.8
1.0